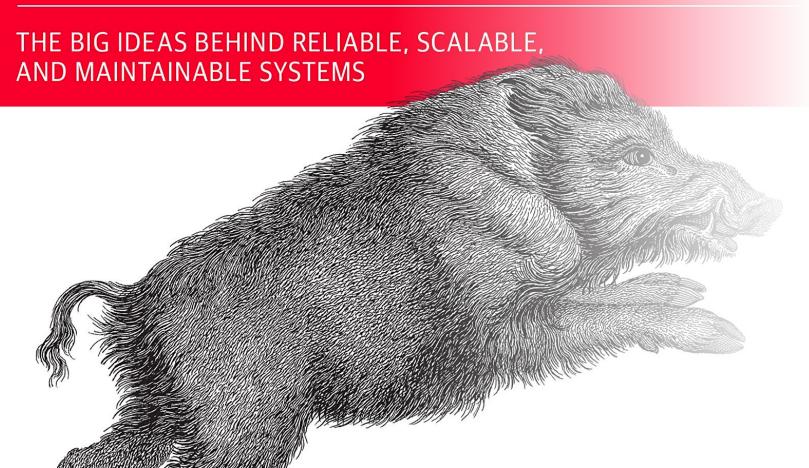
# Data-Intensive – Applications



Chapter 5 : Replication

Why you want to store data across multiple machines?

### Scalability:

If your data volume, read load or write load grows bigger than a single machine can handle

Fault tolerance/High Availability:
 If one machine crashes or some failure occurs, you can enable redundancy

### • Latency:

Users are spread across different locations, serve from location that is geographically close to the users

## How data is distributed across multiple machines?

Replication

**Partitioning** 

### Partition 1, Replica 1

837 ->
Welcom
e Catie

847 ->
Learn
with me
Mate

Mate

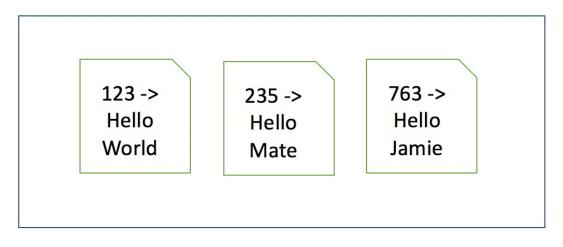
#### Partition 1, Replica 2

837 ->
Welcom
e Catie

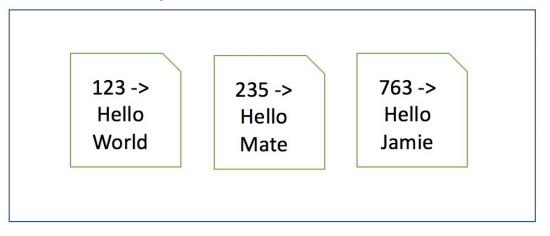
847 ->
Learn
with me
Mate

983 ->
Happy
coding

### Partition 2, Replica 1



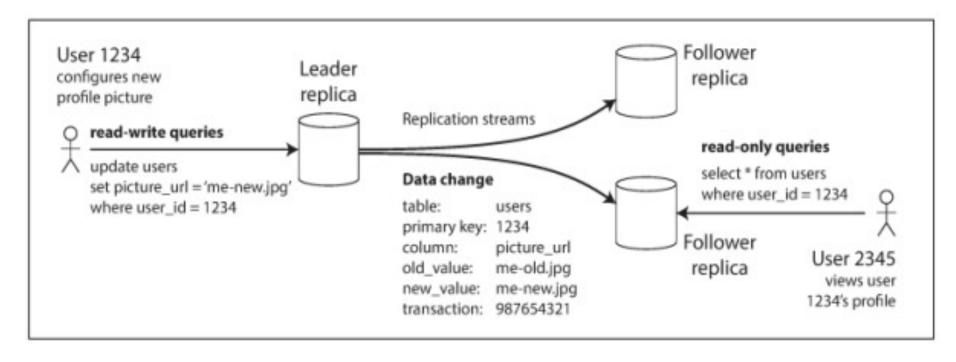
#### Partition 2, Replica 2



## Replication

- Keep data geographically close to users -> Latency
- Fault tolerance Availability
- Scale out for read and write queries
- Problem with replication: Handling changes to replicated data
- Approaches:
  - Single leader
  - Multi leader
  - Leaderless

### Leaders and Followers



Leader based replication, active/passive or master-slave replication

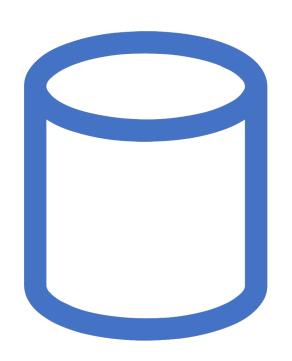
### Synchronous vs Asynchronous Replication

### Synchronous

- Follower is guaranteed to have an up-to-date copy of data that is consistent with the leader.
- If leader suddenly fails, data will be still available on follower.
- If sync follower does not respond, leader must block all writes, till the sync replica is available.
- Generally sync replication in database means- one follower is sync and other followers are async Semi synchronous

### Asynchronous

• If leader fails and is not recoverable, writes get lost, but it guarantees that leader can continue processing writes.



## How can we set up the followers?

- Take consistent snapshot of leader's database at some point in time.
- Copy the snapshot to the new follower node.
- Follower connects to the leader and requests all the data changes that have happened since the snapshot was taken. Mysql – binlog coordinates..
- Now when follower has caught up, it can process the data changes.

How can we achieve high availability using leader-based

User 1234

configures new

profile picture

read-write gueries

set picture\_url = 'me-new.jpg

where user\_id = 1234

update users

Leader

replica

Replication streams

primary key: 1234

Data change

Follower

replica

replica

read-only queries

select \* from users

where user id = 1234

User 2345

1234's profile

replication?

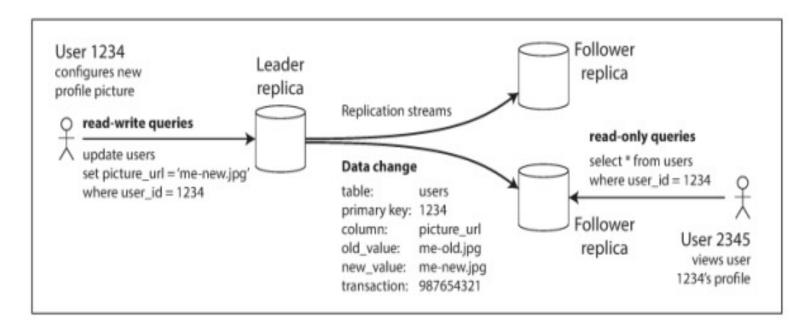
Follower failure – Catch up Recovery

Local disk – contains log

#### Leader failure –Failover

- One of the follower needs to be promoted to the new leader
- Client needs to be reconfigured to send writes to the new leader
- Other followers start consuming data changes from new leader
- Automatic failover process:
  - 1. Determining that leader has failed using timeouts
  - 2. Choosing a new leader Consensus problem
  - 3. Reconfiguring the system to use the new leader

### Issues with automatic failover



- If async replication new leader might not be up-to date. If former leader joins back the cluster, then conflicting writes
- Discarding writes is dangerous.
- Split brain

## Implementation of Replication logs

Statement based replication logs

Write-ahead log(WAL)

Logical(rowbased) log replication

Trigger-based replication

# Statement based replication

 Leader logs every request that it executes and sends that statement to its followers.

### Issues:

- Non-deterministic functions like now(), rand() can generate different results on different replicas.
- If statements are using auto increment column, then multiple concurrent statements can be an issue.
- Statements that trigger procedures etc. can result in different outcomes.

## Write-ahead Logs

- In case of log-structured engines, this WAL is the main place for storage.
- In case of B-trees, which overwrites individual disk blocks, every modification is written to WAL first.
- Log describes data on very low level, it contains details of which bytes were changed in which disk blocks

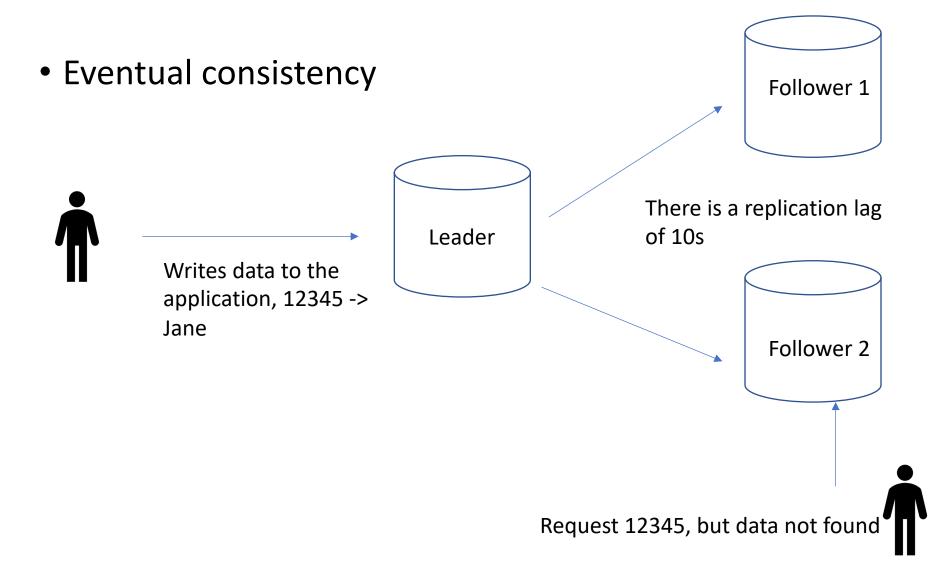
# Logical(row-based) log replication

- Different log formats for replication and storage engines – logical log- distinguished from storage engine's (physical) data representation.
- Logical log contains:
  - Inserted row log contains new values of all columns
  - Deleted row Primary key, and old value of all columns
  - Updated row Primary key and new values of all columns
  - Then a log indicating transaction was committed.

Triggerbased replication

Trigger – code that is automatically executed when a data change occurs in database system.

## Problems with replication lag



# Approaches to Replication Lag

- Reading your own writes
- Read-after-write consistency/ read-yourwrites consistency
- Monotonic reads
- Consistent Prefix reads

### Summary



What is Replication?



Why Replication is needed?



**Single Leader Replication** 



Replication Lag

