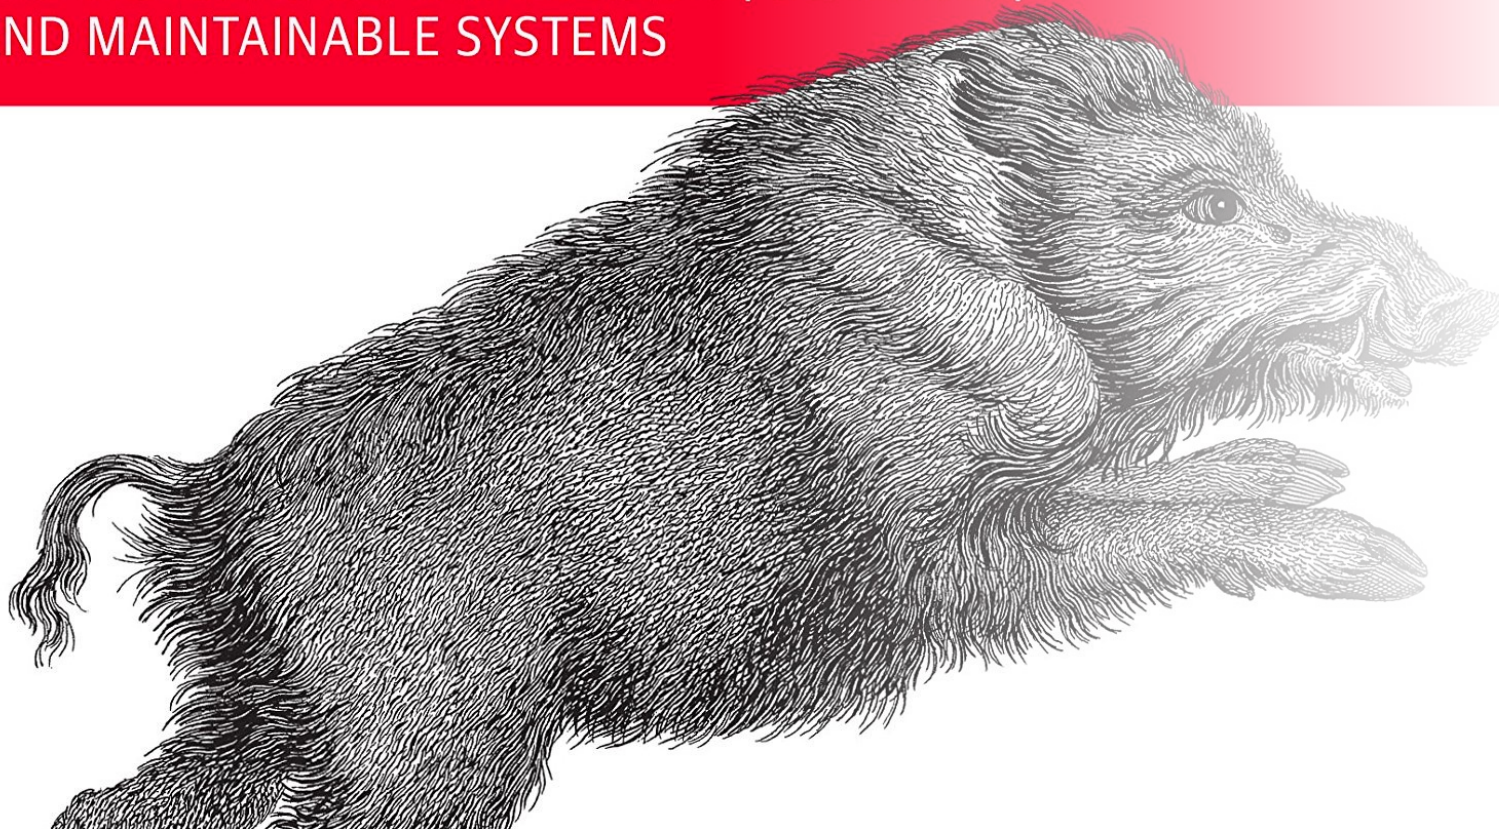


Data-Intensive — Applications

THE BIG IDEAS BEHIND RELIABLE, SCALABLE,
AND MAINTAINABLE SYSTEMS



Chapter 7 : Transactions

What can go wrong in data systems?

- Database software or hardware may fail at any time.
- Application can crash at any time.
- Interruptions in the network.
- Concurrent clients.
- Data partially updated
- Race conditions can cause bugs.

ACID

- **Atomicity :**

Describes what happens if a fault occurs after some of the writes have been processed. Then those writes are grouped together into an atomic transaction and the transaction cannot be completed due to a fault, then the transaction is aborted and database must discard or undo any writes it has made so far in the transaction.

- **Consistency :**

Application specific notion of database being in a good state.

- **Isolation :**

Concurrently executing transactions are isolated from each other, which means that each transaction can pretend that it is the only transaction running in the entire database.

- **Durability :**


Promise that once transaction has been committed successfully, any data it has written will not be forgotten, even if there is a hardware fault or the database crashes.



Single-Object and Multi-Object operations

- Atomicity
 - Isolation
 - BEGIN TRANSACTION and COMMIT statement.
-

Single-object writes



Network connection is interrupted.



If power fails in middle of overwriting the previous value on disk.



If client sees any partially updated value?



Atomicity :
Implemented using a
log for crash recovery



Isolation : Using lock
on each object

Why do we need multi-object transactions?

- In relational data model, a row in one table has a foreign key reference to a row in another table. It ensures that these references remain valid.
- Document data model lacking joins encourages denormalization. When denormalized information needs to be updated, we need to update several documents in one go.
- In databases with secondary indexes, the indexes need to be updated every time you change a value.



Handling errors and aborts

- If transaction actually succeeded, but network failed.
- If error is due to overload, retrying the transaction will make the problem worse not better.
- It is worth retrying only after a transient error but not a permanent error.

Weak Isolation Levels

- Concurrency issues comes when one transaction reads data that is concurrently modified by another transaction or when two transactions try to simultaneously modify the same data.
- Transaction isolation
- Serializable isolation – database guarantees that transactions have the same effect as if they ran serially. It has a performance cost.
- Non serializable isolation levels

Read Committed

- Two guarantees:
- When reading from the database, you will only see data that has been committed(no dirty reads).
- When writing to the database, you will only overwrite data that has been committed.(no dirty writes).

No dirty reads

- A transaction that reads the data from another transaction which is still in uncommitted state.
- Any writes by a transaction only become visible to others when that transaction commits.
- Why dirty reads are prevented?
 - If transaction needs to update several objects, a dirty read means that another transaction may see some of the updates but not others. Data will be seen in partially updated state which is confusing.
 - If a transaction aborts, any writes that has been made need to be rolled back. If db allows dirty reads, that means a transaction may see data that is later rolled back, which is never actually committed to the database.

No dirty writes

- We don't know in which order writes happen when two transactions try to concurrently update the same object in a database. If an earlier writes is a part of transaction that has not yet been committed and later writes overwrite an uncommitted value.
- Why dirty write is needed?
 - If transactions update multiple objects, dirty writes can lead to a bad outcome.

Implementing read committed

- Row level locks – holds lock until that transaction is committed or aborted.
- One long running write transaction can force many read-only transactions to wait until the long running transaction has completed.
- Response time is slow and bad for operability.
- Most databases use dirty reads – for every object that is written, the database remembers both the old committed value and new value set by transaction that holds the write lock.
- When transaction is ongoing, any other transactions that read objects are given old value, until the transaction holding the write lock commits the new value.

Summary



ACID



Single object writes



Multi-object writes



Weak Isolation levels



Read committed Isolation level



Dirty reads



Dirty writes



Thank You!